

How to use this document: Read Section 1 first for the engineering context. Section 2 gives you a step-by-step PINN build guide. Section 3 explains the digital twin workflow. Section 4 is your ready-made business case for management presentation.

Contents

Section 1 — The Engineering Problem

What thermal management means in an exhaust system and why it matters.

Section 2 — PINN Build Guide

Step-by-step instructions: energy equation, training, validation, and deployment.

Section 3 — Digital Twin Workflow

How sensor data flows into the PINN, what outputs are produced, and how systems respond.

Section 4 — Business Case for Management

Value proposition, ROI framing, KPIs, risks, and a ready-made presentation structure.

Section 1 — The Engineering Problem

What Is Exhaust Thermal Management?

Exhaust thermal management is the engineering discipline of controlling temperature distribution across the exhaust system to achieve three goals simultaneously: catalytic converter light-off (reaching the minimum temperature for pollutant conversion), component protection (preventing metal fatigue from thermal extremes), and system efficiency (minimising heat loss that reduces catalyst effectiveness).

Why Temperature Distribution Matters

A catalytic converter does not work uniformly. Exhaust gas enters the catalyst face with a velocity profile shaped by the upstream pipe geometry — bends, junctions, and area changes create regions of high and low velocity. This means different zones of the catalyst heat up at different rates. The zones that receive low flow heat slowly and stay below light-off temperature for longer, producing excess emissions. The zones with high flow may overheat and degrade the catalyst washcoat.

Critical Threshold: Catalyst light-off occurs at approximately 250–300°C. Below this temperature, conversion efficiency is near zero. During a cold start, the first 60–120 seconds produce the majority of a vehicle's total regulated emissions.

The Gap Traditional Methods Leave

Traditional Approach — Limitations	PINN Digital Twin — Solution
✗ CFD simulations take 4–24 hours — no real-time updates	✓ PINN inference returns full field in milliseconds
✗ Thermocouples only measure at fixed sensor locations	✓ Every mesh node predicted — no blind spots
✗ Lookup tables are static — cannot track transients	✓ Live boundary conditions track transients continuously
✗ Cold-start events last 60–120 seconds — too fast for CFD	✓ Cold-start evolution predicted from first engine cycle
✗ Hot spots between sensors are completely invisible	✓ Complete 3D spatial map at all times
✗ Every new geometry needs a full new test campaign	✓ Retrain once — cover all operating conditions

What Success Looks Like

<10 ms

PINN inference time

100%

Surface coverage vs point
sensors

60–90 s

Cold-start window fully
tracked

Zero

Extra hardware sensors
needed

Section 2 — PINN Build Guide

Starting Point: Your current code already solves Navier-Stokes for velocity (u , v , w) and pressure (p). Thermal management requires adding temperature (T) as a 5th output variable by incorporating the energy equation. Everything else — geometry, mesh, training loop, and solver — stays the same.

Step-by-Step PINN Build Instructions

1

Add the Energy Equation to Your Physics

The energy equation describes how heat is transported through the fluid by convection and diffusion. In your current code, add the AdvectionDiffusion PDE from PhysicsNeMo alongside the existing Navier-Stokes equations.

In your `run()` function, after the NavierStokes instantiation, add:

```
from physicsnemo.sym.eq.pdes.advection_diffusion import AdvectionDiffusion
ad = AdvectionDiffusion(T='T', rho=1.0, D=k_eff, dim=3, time=False)
```

Where k_{eff} is the effective thermal diffusivity of exhaust gas ($\sim 2.5 \times 10^{-5} \text{ m}^2/\text{s}$ at 500°C).

Add `ad.make_nodes()` to your nodes list alongside the existing `ns.make_nodes()`.

2

Update the Network Architecture

Your existing `flow_net` predicts $[u, v, w, p]$. You need to add T as a 5th output. In your config (`conf/config.yaml`), add T to the output keys:

```
output_keys: [Key('u'), Key('v'), Key('w'), Key('p'), Key('T')]
```

The fully connected architecture will handle the 5th output automatically. No other architecture changes are required. The network will learn u, v, w, p, T simultaneously — the physics equations couple them during training.

3

Define Thermal Boundary Conditions

You need to specify temperature at every boundary where you already have velocity BCs:

INLET: $T = T_{\text{inlet}}$ (measured upstream temperature, e.g. 600°C cold start, 750°C hot)

OUTLET: Zero gradient (Neumann BC — no temperature constraint needed)

WALL: $T = T_{\text{wall}}$ (measured wall temperature or adiabatic: $dT/dn = 0$)

Update your inlet `PointwiseBoundaryConstraint` to include 'T': `T_inlet_value`

Update your noslip `PointwiseBoundaryConstraint` to include 'T': `T_wall_value`

For adiabatic walls, use `NeumannBC` with 'normal_gradient_T': 0 instead.

4

Train Across a Range of Operating Conditions

For digital twin deployment, the PINN must work across all engine operating points, not just a single steady state. Use parametric training by making T_{inlet} a Symbol:

```
T_in = Symbol('T_in')
```

```
inlet BC: outvar={'u': u_profile, 'v': v_profile, 'w': w_profile, 'T': T_in}
```

```
param_ranges={T_in: (400, 1100)} # Kelvin range for exhaust inlet
```

Similarly parameterise inlet velocity to cover the full RPM/load range.

This trains a single network that acts as a lookup table for ALL conditions at once.

5

Add Validation Data

Validation is critical before deployment. You need reference data to confirm accuracy:

Option A (Best): Run 3–5 CHT CFD simulations at representative operating points. Export temperature fields as CSV. Use PointwiseValidator to compare PINN vs CFD.

Option B: Use thermocouple measurements from dyno testing as ground truth. Map sensor locations to mesh coordinates, normalise the same way as your geometry.

Target accuracy: Temperature error < 5% across the operating range, especially within the catalyst zone (most safety-critical region).

6

Export and Wrap for Real-Time Inference

After training, export the model and wrap it in a Python inference function:

```
def predict_thermal(T_in, vel_in): # T in Kelvin, vel in m/s
```

```
    invar = {'x': mesh_x, 'y': mesh_y, 'z': mesh_z,
```

```
            'T_in': np.full_like(mesh_x, T_in),
```

```
            'vel_in': np.full_like(mesh_x, vel_in)}
```

```
    return model.predict(invar) # returns u,v,w,p,T at all nodes
```

This single function call is all the digital twin needs. It runs in ~10 ms.

Training Time Estimate: Parametric training with temperature typically converges in 24–72 hours on a single GPU (V100 or better). Start with 100,000 iterations and monitor the energy equation residual alongside momentum residuals in TensorBoard.

Checklist: Is Your PINN Ready for Deployment?


☐	Energy equation residual converges to < 1e-3
☐	Temperature field validated against at least 3 CFD reference points
☐	Cold-start condition (low T_{in}) tested — no non-physical temperatures
☐	Hot operating condition (high T_{in}) tested — no numerical overflow

■	Inference time benchmarked: target < 50 ms on CPU, < 10 ms on GPU
■	Catalyst zone nodes identified and labelled in mesh for fast extraction
■	Model saved with correct input normalisation parameters

Section 3 — Digital Twin Workflow

Complete System Architecture

The digital twin sits between the physical sensors and the response systems. It receives real-world data every engine cycle, runs PINN inference, and delivers actionable outputs to three downstream consumers: the ECU for active control, the monitoring dashboard for engineering visibility, and the data logger for long-term thermal history analysis.



Detailed Workflow — Step by Step

#	Step	What Happens	Time	Output
1	Engine cycle begins	ECU fires. RPM, load, and injection timing are logged.	0 ms	Operating point definition

#	Step	What Happens	Time	Output
2	Sensor data acquired	Upstream thermocouple reads T_in. MAF sensor reads mass flow rate.	5–10 ms	T_in (°C), mdot (kg/s)
3	Boundary conditions formed	T_in and velocity magnitude converted to PINN input format. Normalised to training scale.	1 ms	Input tensor for PINN
4	PINN forward pass	Neural network evaluates u, v, w, p, T at all mesh nodes simultaneously.	5–15 ms	Full 3D field: 5 × N values
5	Catalyst zone extraction	Temperature values at catalyst mesh nodes extracted. Min, mean, max computed.	1 ms	T_cat_min, T_cat_mean, T_cat_max
6	Light-off decision	If T_cat_mean > 280°C: flag ACTIVE. If < 200°C: flag COLD. Between: WARMING.	< 1 ms	Status flag + time-to-light-off estimate
7	Hot-spot detection	Surface nodes with T > T_threshold_material flagged. Location coordinates logged.	1 ms	Hot-spot list with coordinates and severity
8	ECU command generation	If cold: request fuel enrichment to raise exhaust temperature. If hot-spot: request EGR.	1 ms	CAN bus message to ECU
9	Dashboard update	New temperature field rendered as 3D heatmap. Gauge values updated.	Async	Live visualisation for engineer
10	Data logged	Full field snapshot stored at 1 Hz. Summary metrics stored at 10 Hz.	Async	Database entry for trend analysis

Total Loop Time: Steps 1–8 complete in under 30 ms, well within a 100 ms ECU cycle. Steps 9–10 run asynchronously and do not block the control loop.

What the Dashboard Shows

■ 3D Temperature Heatmap

Colour-mapped visualisation of T across the full exhaust geometry. Updated at 1 Hz. Engineer can rotate, zoom, and inspect any node. Red zones indicate hot spots above threshold.

■ Catalyst Light-Off Gauge

Large indicator showing COLD / WARMING / ACTIVE status with current mean catalyst temperature. Timer showing seconds since engine start and estimated seconds to light-off completion.

■ Hot Spot Alert Panel

List of nodes exceeding material temperature limits, with location description (e.g. 'Bend 2 — outer wall') and current temperature value. Alerts ranked by severity.

■ Thermal History Chart

Time-series of T_cat_mean from engine start. Annotates cold-start window, light-off event, and any thermal exceedance events. Exported per drive cycle.

■ ECU Command Log

Record of all enrichment or EGR commands sent with timestamp and reason. Allows engineer to audit control actions and tune thresholds.

Section 4 — Business Case for Management

How to use this section: Present slides in this order — Problem, Scale of Impact, Our Solution, How It Works (one diagram), Evidence It Works, Investment Required, Timeline, Return on Investment. Each sub-section below maps to one slide.

Slide 1 — The Problem (30 seconds to explain)

Say this to your manager: 'We are currently flying blind on exhaust temperatures. During the most critical window — cold start — we have no real-time view of what is happening inside the catalyst. We either run expensive simulations that take hours, or we rely on a handful of point sensors that miss the spatial picture entirely. The result is either excess emissions we cannot see, or thermal damage we detect too late.'

Slide 2 — Scale of Impact (Use These Numbers)

Impact Area	Current Situation	Business Cost
Cold-Start Emissions	60–70% of total regulated emissions occur in first 2 minutes	Regulatory non-compliance risk, type approval failure
Thermal Damage Discovery	Detected only after failure — no predictive signal	Unplanned warranty claims, late-stage design changes
Simulation Bottleneck	4–24 hours per operating point for CHT CFD	Delays thermal validation by weeks per programme
Sensor Infrastructure	Dense thermocouple arrays needed per variant	Hardware cost + installation time per test vehicle

Slide 3 — Our Solution in One Sentence

"We train a physics-informed neural network on our exhaust geometry once, then use it as a real-time digital twin that predicts full 3D temperature distribution in under 10 milliseconds from cheap upstream sensors — no additional hardware required."

Follow this with the three-line explainer: (1) It learns the physics of our geometry from simulation data. (2) At runtime it takes sensor readings and produces the complete thermal state instantly. (3) The result feeds directly into the ECU for active control.

Slide 4 — Evidence It Works

Use these reference points to build credibility with technical managers:

- NVIDIA PhysicsNeMo (the framework used) is deployed in production for CFD acceleration at BMW, Ford, and multiple Tier-1 suppliers
- Published PINN literature demonstrates temperature field prediction accuracy within 3–5% of full CHT CFD on similar exhaust geometries
- Our existing code already solves Navier-Stokes on this exact geometry — adding temperature is a targeted extension, not a restart
- Inference time is deterministic: the network is a fixed matrix multiplication chain. It will always run in the same time window
- The approach is geometry-specific — trained on our exhaust, not a generic model. This is an asset, not a limitation

Slide 5 — Investment Required

Resource	What Is Needed	Estimate
Engineering Time	1 engineer to extend PINN code and build inference wrapper	3–4 weeks
Compute (Training)	GPU server for 24–72 hour training run (AWS p3.2xlarge or equivalent)	\$200–500 one-time
Validation Data	3–5 reference CHT CFD simulations at key operating points	Use existing CFD pipeline
Dashboard	Simple Python/web frontend to visualise temperature field	1 week additional
Ongoing Cost	Inference cost is near zero — runs on existing ECU or edge hardware	Minimal

Slide 6 — Delivery Timeline

Week	Milestone	Deliverable
Week 1–2	Add energy equation, update architecture, first training run	PINN with temperature output — first results
Week 3	Validate against CFD reference data, tune boundary conditions	Validation report: PINN vs CFD accuracy table
Week 4	Build inference wrapper and sensor integration layer	Live inference demo from simulated sensor feed
Week 5	Dashboard development and ECU command interface	Working dashboard with real-time 3D temperature display
Week 6	Integration testing, threshold tuning, manager demo	Go/No-go decision on full deployment

Slide 7 — Return on Investment

Frame the ROI around four measurable outcomes. Present these as before/after comparisons:

■ Emissions Compliance Risk

BEFORE: Cold-start emissions spike is invisible until physical test failure.

AFTER: Real-time light-off prediction enables proactive calibration intervention. Compliance verified on every drive cycle, not just scheduled test events.

■ Thermal Damage Prevention

BEFORE: Hot spots detected only when physical components fail during endurance testing.

AFTER: Hot-spot severity known in real time. Design changes made before components are built.

■ Simulation Cost Reduction

BEFORE: 3–5 CHT CFD runs per geometry variant at 4–24 hours each for thermal validation.

AFTER: PINN covers all conditions within trained range. CFD reserved for training data only.

■ Development Programme Acceleration

BEFORE: Thermal validation causes 2–3 week delays while CFD queue clears.

AFTER: Thermal state predicted instantly for any operating condition — validation is always available.

How to Close the Presentation: 'We are not proposing to replace CFD. We are proposing to make our existing CFD knowledge available instantly, at any operating point, continuously — without running a single new simulation. The investment is 5 weeks of engineering time. The return is a permanent real-time thermal intelligence capability on every vehicle that runs our exhaust geometry.'

KPIs to Track After Deployment

KPI	Definition	Target
Light-Off Time	Seconds from cold start to catalyst ACTIVE status	Reduction vs baseline > 10%
Prediction Accuracy	PINN T vs reference thermocouple at 3 validation points	< 5% mean absolute error
System Latency	Time from sensor read to ECU command output	< 30 ms total
Thermal Exceedance Events	Number of cycles where T _{wall} > material limit	100% detected before component damage
CFD Hours Saved	Reduction in CHT CFD runs for thermal validation per programme	> 50% reduction